

Project Progress Report

- Presented by Atharv Kumar and Yashasvi

Overview:

We have made substantial progress in understanding and experimenting with **font style generation and interpolation**. Our work to date can be summarized in four major contributions:

1. In-depth Analysis of the WordStylist Paper

- Studied the methodology, architecture, training strategy, and evaluation metrics (FID, CER, WER, writer classification, and retrieval) and Understood how **Latent Diffusion Models (LDMs)** outperform GANs in generating high-quality, style-consistent handwriting.

2. Survey of Existing Font Interpolation Techniques

- Explored GAN-based methods (GANwriting, SmartPatch, ScrabbleGAN), Transformer models, and style-transfer approaches.


3. Study and Initial Implementation of Bezier Curve-Based Style Interpolation

- Investigated bezier curve-based techniques for **smooth style blending**. Attempted to code and implement curve interpolation between latent style embeddings to morph handwriting styles.

4. Metric-Based Evaluation Understanding

- Learned how to quantitatively assess synthetic handwriting quality using FID, CER/WER, writer classification accuracy, and retrieval performance.

In-depth Analysis of the WordStylist Paper:



(The overall architecture)

Overview of Diagram Sections

Right Side(Text and style conditioning), Middle(Latent diffusion process (UNet)),Left Side (Image encoder/decoder (VAE))

1. Text & Style Conditioning Module (Right Side)

Input Text (τ): "that"

Black solid arrows → Training phase

Red dashed arrows → Sampling (generation) phase

→ Text Conditioning Pipeline:

Embedding: Each character is converted into a numerical vector.

Positional Encoding: Adds position information to the character embeddings.

Self-Attention: Captures dependencies and context between characters.

Output: Generates a final text condition vector c_τ , which represents the content of the word.

→ Style Condition: Y

Writer style is input as a **style class index Y** , embedded into a vector.

2. Latent Diffusion Process (Middle Section)

→ Forward Process (Training):

The real image of "that" is encoded by the VAE encoder **V_E** into a latent vector **z** and here noise is **gradually added** over time steps: $\mathbf{z} \rightarrow \mathbf{z}_1 \rightarrow \mathbf{z}_2 \rightarrow \dots \rightarrow \mathbf{z}_T$.

→ UNet (Core Denoising Network):

At any timestep **t**, the noisy latent \mathbf{z}_t , the text condition \mathbf{c}_T , and the style condition \mathbf{Y} are fed into the UNet. The UNet (with **attention blocks: Q/K/V**) **predicts the noise** in \mathbf{z}_t , to help reverse it.

During training, the UNet is optimized to predict the actual noise using L2 loss.

3. Reverse Sampling (Left Direction - Red Arrows)

→ Sampling Pipeline:

Start with a random Gaussian latent vector \mathbf{z}_T . Use the trained UNet to progressively denoise the latent vector: $\mathbf{z}_T \rightarrow \mathbf{z}_{T-1} \rightarrow \dots \rightarrow \mathbf{z}_0$.

The final clean latent \mathbf{z}_0 is passed through the VAE decoder **V_D**, which generates a synthetic handwriting image of the input word "that" in the chosen writer style.

How the Algorithm Works — Step-by-Step:

A. Forward (Training) Process

1. **Encoding Real Image to Latent Space:**
 - Real handwritten word images are passed through a **Variational Autoencoder (VAE) encoder** to get a **latent representation \mathbf{z}** .
2. **Adding Noise in Timesteps:**
 - A **Gaussian noise** is gradually added to \mathbf{z} over **T = 1000 timesteps** using a noise schedule. Here, β_t increases linearly from 1e-4 to 0.02.

$$q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t} x_{t-1}, \beta_t I)$$

3. Conditioning Inputs:

- **Text condition \mathbf{T} :** Passed through a **Content Encoder (CE)** — a Transformer-based module with **positional encoding** and **self-attention**.
- **Style condition \mathbf{Y} :** Writer ID is embedded and added to the timestep embedding.

3. Noise Prediction:

- A **U-Net** architecture is used to predict noise at each timestep using the noisy latent \mathbf{z}_t , style, and text condition.

4. Loss Function:

- The network minimizes the **L2 loss** between true noise ϵ and predicted noise $\hat{\epsilon}$.

$$L = \mathbb{E}_{x_0, t, \epsilon} [\|\epsilon - \hat{\epsilon}_\theta(x_t, t)\|^2]$$

B. Reverse (Sampling) Process

1. Sampling from Pure Noise:

- Begin with $\mathbf{z}_T \sim \mathcal{N}(0, I)$ (pure noise).

2. Reverse Denoising Process:

- Use the trained model to iteratively **denoise** from \mathbf{z}_T to \mathbf{z}_0 :

$$p_\theta(x_{t-1} | x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

-

3. Decoding to Pixel Space:

- The denoised latent `z_0` is passed through a **VAE decoder** to produce the final **synthetic word image**.

4. Sampling Optimization:

- Authors reduce timesteps from 1000 to **600**, balancing quality and speed (~12 sec/image).

3. Training Strategy

- **Dataset:** IAM Offline Handwriting Dataset (word-level, 339 writers).
- **Preprocessing:** Words with 2–7 characters, images resized to $64 \times W$, padded if necessary.
- **Training:**
 - Epochs: 1000
 - Batch size: 224
 - Optimizer: AdamW, LR = $1e-4$

4. Evaluation Metrics

A. Image Quality

- **Fréchet Inception Distance (FID):**
 - Lower is better.
 - WordStylist: **22.74**
 - SmartPatch: 22.55
 - GANwriting: 29.94

B. Text Recognition (HTR) Performance

- Model: CNN-LSTM with CTC loss
- Metrics:
 - **Character Error Rate (CER)**
 - **Word Error Rate (WER)**
- Results (training on synthetic + real data):

Training Data	CER (%)	WER (%)
Real IAM only	4.86	14.11
WordStylist only	8.80	21.93
Real IAM + WordStylist	4.67	13.28

C. Writer Style Preservation

- **Writer Classification Accuracy (ResNet18):**
 - WordStylist: **70.67%**
 - GANwriting / SmartPatch: ~5%
- **Writer Retrieval (mAP & Top-1 Accuracy):**
 - mAP (Mean Average Precision):
 - WordStylist: **97.84%**
 - Real IAM: 97.61%
 - GANwriting / SmartPatch: ~7%
- **t-SNE plots** also show WordStylist clusters match real handwriting better than GANs.

Study and Initial Implementation of Bezier Curve-Based Style Interpolation :

1. Style Embedding Extraction

Extract **latent style vectors** (e.g., Y_1 , Y_2 , Y_3) from a trained model like WordStylist. Each vector represents a unique handwriting style.

2. Understand Bezier Curve Basics

Use control points to define a smooth curve in latent space. here we can see this in example for cubic:

$$B(t) = (1-t)^3 Y_0 + 3(1-t)^2 t Y_1 + 3(1-t)t^2 Y_2 + t^3 Y_3$$

where $t \in [0, 1]$.

3. Implement Curve Interpolation

Wrote code to **compute Bezier curve points** between style embeddings using NumPy or PyTorch.

Sampled multiple values of t (e.g., 0.0 to 1.0) to generate intermediate style vectors along the curve.

4. Generate Interpolated Handwriting

Used these **interpolated style vectors** in place of fixed style inputs in the generative model. Kept the input word fixed (e.g., "hello") to visualize how the **style morphs gradually**.

5. Observe and Analyze Results

Verified that the output images showed **smooth style transitions**.

Compared with linear interpolation and found Bezier-based blending more **fluid and natural**.

Survey of Existing Font Interpolation Techniques:

As part of our project, we conducted a thorough survey of existing methodologies used for font or handwriting style generation and interpolation. The goal was to understand the state-of-the-art techniques, their architectures, and how they manage style control, text conditioning, and generalization to unseen data.

1. GAN-Based Methodologies

a. GANwriting

- Uses a style vector (writer identity) and text embedding as inputs to generate word-level handwritten text. One of the strengths is that it is able to generalize to out-of-vocabulary (OOV) words and produces visually realistic images with decent text alignment.
Weakness include style preservation is inconsistent and lacks flexibility in handling unseen styles (few-shot generalization) and training is unstable due to adversarial losses.

b. SmartPatch

- Improvement Over GANwriting: Introduces a patch-based discriminator to reduce local artifacts in generated images.
- Strengths: Smoother and cleaner outputs than GANwriting with Better structural consistency in characters.
- Weaknesses: Still struggles with global style coherence and Requires large labeled data for effective training.

c. ScrabbleGAN

- Designed to generate longer handwritten sequences (multi-word or sentence-level) It is a semi-supervised approach that doesn't require full style labels.
- Strengths: Can synthesize text of varying lengths and reduces annotation overhead.
- Weaknesses: Weaker control over style consistency and slightly lower image quality compared to fully supervised models.

2. Transformer-Based Approaches

- The Architecture includes Encoder-decoder with self-attention mechanisms. The input is the Style features from CNN + textual content. Whereas the output is Handwritten sequence images.
- Strengths: Better handling of long dependencies (useful for longer words or sentences) and Flexibility in integrating multiple modalities (e.g., style + text).
- Weaknesses: Requires large datasets and heavy computation, Overfitting risk in low-resource settings and Slower inference compared to CNN-based GANs.

3. Style-Transfer Approaches

These models aim to transfer handwriting style from one domain to another. It basically involves CycleGAN-Based Methods where the Style transfer for historical document generation.

Process involves Converting modern documents (e.g., LaTeX) into stylized handwritten versions using unpaired image translation.

- Strengths: Works well even without paired training data and Good for document-level style conversion.
- Weaknesses: Cannot generate new words or sentences and Hard to control specific text content.

Comparative Summary

Method	Text Control	Style Control	Generalization	Quality	Complexity
GANwriting	✓ Yes	⚠ Partial	✗ Weak	✓ Good	⚠ Moderate
SmartPatch	✓ Yes	⚠ Partial	✗ Weak	✓ Better	⚠ High
ScrabbleGAN	✓ Yes	⚠ Weak	✓ OK	⚠ Average	✓ Scalable
Transformers	✓ Yes	✓ Strong	✓ Moderate	✓ High	✗ Heavy
CycleGAN	✗ No	✓ Style-only	⚠ Limited	✓ Good	⚠ Moderate

Future Work:

Going forward, our project will focus on building a robust **font interpolation pipeline** using **Bezier curve-based blending** within **Diffusion** and **GAN frameworks**. The aim is to generate smooth, style-consistent handwritten text across multiple fonts and languages, including Indic scripts. Key future tasks include:

- Implement Bezier curve-based interpolation between font/style embeddings.
- Benchmark Diffusion vs GAN models on quality, style fidelity, and text recognition.